

David Südholt

Joshua Reiss

Machine learning of physical models for voice synthesis

In the task of sound matching simple physical models of voice production to real human recordings, white-box optimization with gradient descent and differentiable operations outperforms black-box parameter prediction.

With the goal of matching a digital waveguide to a target pitch, can the length of a fractional delay line be estimated using differentiable allpass filters?

Tongue
Parameters,
Constrictions

Base
Diameter